
Forkcast

The trust layer for restaurant nutrition decisions

Business Plan — Version 2.0 · July 2026

Prepared for pre-seed conversations.

THE PRODUCT IN ONE SENTENCE

Forkcast tells diners what to order *before* they order — scored against their own body and goals — then closes the loop: order, confirm, log, measure. It is the only product connecting pre-order decision support, verified independent-restaurant nutrition data, and post-order confirmed logging.

SOURCING & INTEGRITY

All market figures are sourced (CDC, USDA ERS, AHRQ, peer-reviewed literature, company filings/press; see the accompanying Market Research Dossier). Competitive claims come from live site captures dated July 15, 2026 (Competitive Scan appendix). Nothing in this document claims users, partners, revenue, or outcomes that do not exist.

A complete working product accompanies this plan; a live demonstration is available on request.

Appendices: Financial Model · Market Research Dossier · Menu Collection Protocol · System Architecture

0 The answer first

Americans eat out more than ever, decide blind, and every existing app reacts only after the plate is empty. The category leader for “eat out healthy” markets AI guesses across 22 million locations as facts; the logging giants own the aftermath; the delivery marketplaces own the transaction. Nobody owns the *decision* — and nobody at all serves the roughly 65% of U.S. restaurants that are independent and exempt from federal menu labeling.

Forkcast owns the decision, and earns the right to with honesty no competitor structurally can match: every nutrition number carries its source and a $\pm$ range; restaurants review and correct their own data through a versioned public process; sponsored placement can never touch the score; and the personal calorie target itself recalibrates from the user’s own measured energy balance. The full product — decision → order → restaurant terminal → confirmed log → adaptive measurement — **exists and works today.**

The ask: ~\$1.0–1.5M pre-seed SAFE (~\$5–6M post) plus non-dilutive SBIR, to run the Boston pilot against six pre-registered metrics, convert the first independent restaurants into verified partners, and prove retention economics.

1 Situation — eating out is where American diets are decided

Fact	Figure	Source
Food spending eaten away from home	58.5%	USDA ERS, 2024 (record; \$4,485/capita, +12% YoY)
Daily calories from food away from home	~1/3	USDA ERS / NHANES (directional)
Restaurant vs. home meal	+200 kcal, +350 mg Na	USDA ERS
Diners underestimating restaurant calories	2 in 3	USDA/NHANES 2024 (~25% by 500+ kcal)
Behavior change from mandatory calorie labels	~24 kcal/order	Peer-reviewed labeling literature
U.S. adults obese / overweight	40.3% + 31.7%	CDC NCHS, 2024 (~72% combined)
Diet-related disease cost	~\$334B/yr	CDC; AHRQ MEPS (obesity ~\$173B, 2019\$)

2 Complication — three structural failures nobody has fixed

- The data doesn’t exist where people actually eat.** FDA menu labeling (21 CFR 101.11) binds only 20+-location chains. The independent majority publish nothing, and no commercial database covers them.
- The timing is wrong everywhere.** Logging apps (MyFitnessPal: 220M registered, ~30M MAU) act after eating — accounting, not coaching — and 70–80% of calorie-app users quit within two weeks. Labels act at the menu but move behavior ~24 kcal. Delivery apps act at the transaction with zero nutrition intelligence.

3. **The new “AI menu” entrants sell certainty they don’t have.** MenuFit (live capture, July 2026): “#1 App For Eating Out Healthy,” 22.3M locations, meals that “perfectly align with your macros” — no sources, no error ranges, no corrections, no restaurant participation, no web product, and a recommendation that dead-ends. Credibility is the category’s open flank.

3 Resolution — the trust layer, already built

Forkcast turns a health profile (Mifflin-St Jeor BMR → activity-scaled TDEE → goal-adjusted targets; ISSN macros; CDC BMI classes; allergies and conditions) into an explainable per-dish **Fit Score**, ranks every nearby menu around what’s left of the user’s day, and then — unlike every competitor — finishes the job:




Exhibit 1 — The closed loop. Competitors stop at the recommendation (or start after the meal); Forkcast connects the decision to the record. Highlighted steps are where every competitor drops out.

As built — July 2026

Capability (all live)	Why it matters
Fit Score engine + condition/allergen personalization	Explainable, safety-aware ranking; allergen matches excluded from recommendations; condition advisories on every dish
Full ordering loop: basket → checkout → tracking → “Log this meal?”	The decision connects to the record; every logged restaurant meal carries an order reference
Restaurant partner terminal	Order queue with customer allergy flags, prep quotes, and versioned, timestamped menu corrections surfacing publicly on dish pages
Three-tier data provenance	<i>Partner-verified</i> / <i>Restaurant-published</i> (Sweetgreen captured live July 2026; The Halal Guys’ published guide) / <i>Estimated ± 15%</i> (real Boston independents, labeled)
Adaptive metabolic calibration	Target TDEE back-calculated from the user’s own logged intake + weight change, confidence-blended — the state-of-the-art approach, attached to restaurant decisions
Evidence system	Exportable meal log with source/portion/confidence per entry; a public Impact page with six pre-registered pilot metrics and a public source ledger
Role-based accounts; web-first, shareable dish pages	Two-sided from day one; app-only competitors have no web surface at all

Deliberately not built: payment processing, live delivery dispatch, partner handoff links — all shown as clearly-labeled integration states, never simulated as real.

System architecture




Exhibit 2 — Architecture. Three surfaces over shared engines over the data asset. Every arrow that touches the data layer makes it more defensible: corrections from restaurants, order-confirmed ground truth from diners. Full detail in the System Architecture appendix.

4 Market

U.S.-only; reasoning explicit so each number can be challenged (derivations in the Market Research Dossier).

- **TAM ≈ \$11.1B/yr.** ~186M U.S. adults with elevated BMI (≈72% of 258M; CDC 2024) × ~\$60 blended annual ARPU (anchored between MyFitnessPal Premium \$79.99/yr and Cal AI \$29.99/yr). Triangulates with the independently sized calorie-tracking app market (\$5.5–14B, 2025; 12–20% CAGR).
- **SAM ≈ \$3.3B/yr.** ~55.7M goal-oriented active dieters who eat out often (~30% of the pool) × \$60.
- **SOM ≈ \$40–90M ARR (5-yr).** 0.5–2.0% of SAM (0.28–1.1M payers; 1% ≈ \$33M consumer ARR) plus restaurant retail-media and B2B2C legs. **Retention-bound, not demand-bound.**
- **Velocity proof:** Cal AI reached ~15M downloads and \$30M+ ARR in under two years before MyFitnessPal acquired it (Dec 2025).




Exhibit 3 — Market pyramid (widths to scale within each tier's share of the tier above).

5 Competition — and the flank they all leave open

Player	Owns	Structurally absent
MenuFit / PlateMate / Nuuro	“Know what to order”; AI menu scoring; influencer GTM	Sources, error ranges, corrections, restaurant participation, ordering/logging loop, web presence; guilt-based framing excludes the health-condition market
MyFitnessPal	Retrospective logging at scale	Acts after eating; restaurant data weakest (same dish 400–1,200 kcal); 70–80% two-week churn; first revenue decline (–5.7%, 2025)
DoorDash / Uber Eats	Ordering logistics, merchant graph, \$1B+ ads	Zero personal-nutrition intelligence; sponsored placement <i>is</i> the product
Sweetgreen / CAVA	Excellent single-brand transparency	Cannot be brand-agnostic
Foodsmart / Nourish	Payer-funded dietitian care	Care delivery, not the moment of ordering




Exhibit 4 — Positioning. The upper-right quadrant — decide before ordering, on data you can prove — is empty. Chains publish but serve only themselves; the AI entrants act early but cannot prove a number; the giants act late or not at all.

Whitespace in one line: *nobody connects pre-order decision support, verified independent-restaurant data, and post-order confirmed logging — and nobody can prove their numbers.*

The moat compounds through participation: every verified menu, every versioned correction, every order-confirmed meal (ground truth vs. estimate — a validation dataset no retrospective logger can construct), and every calibration cycle makes the data asset harder to replicate than any vision model.

6 Business model — three legs, sequenced

1. **Consumer subscription (launch) — priced to win.** The closed loop is **free forever** (charging for logging is the incumbents’ most-hated move, and free order volume feeds the restaurant leg). Premium — calibration, analytics, unlimited photo AI, condition packs — at **\$4.99/mo or \$39.99/yr**: half of MyFitnessPal’s annual,

- 4–14× under every incumbent monthly. Gate growth on **LTV:CAC $\geq 3:1$** , ~6-month payback, funded by a ~\$0-CAC channel no logger has: restaurant-partner co-marketing.
- Restaurant leg — every dollar earned (incentive ladder).** Rung 1: **Verified Partner is free forever** (terminal, allergy flags, corrections, badge) — their participation improves our data asset. Rung 2: **6% commission on Forkcast-originated pickup orders** — against the 15–30% they already pay marketplaces; they pay only when we deliver a customer. Rung 3: optional **featured placement from \$99/mo**, bought after the order channel proves volume — and **never affecting Fit Scores**. Blended yield ~\$1,050/location/yr.
 - B2B2C covered nutrition (durable).** Payer/employer contracts on covered MNT (CPT 97802–97804); GLP-1 companion programs; Medicaid expansion as optionality.

The price war, won on entry

Product	Monthly	Annual	Note
MyFitnessPal Premium	\$19.99	\$79.99	Premium+ \$24.99/\$99.99; paywall backlash, first revenue decline
Noom / coaching apps	\$40–70	~\$209	\$62M dark-pattern settlement
MacroFactor	\$11.99	\$71.99	Monetizes accuracy/adaptivity — our closest analog
Cal AI	—	\$29.99	Photo-only volume play
Forkcast Premium	\$4.99	\$39.99	The closed loop itself stays free forever

Unit economics at the new price (v2): blended gross ARPU ~\$42/yr; CAC ~\$15 (restaurant co-marketing as a ~\$0-CAC channel); **LTV:CAC $\approx 3.3:1$, payback ≈ 6 months**. The honest trade: Year-5 base ARR moves from \$58M (v1) to a more defensible **\$37M** — low end of the \$30–90M SOM envelope — with named upside paths (featured adoption, order volume, earned ARPU lifts, GLP-1 employer channel) back toward \$50M+. Pricing headroom is *upward*: verified trust can raise prices; eroded trust can only cut. Full build in the Financial Model appendix.

7 Go-to-market — density, then proof, then scale

Phase 1 (months 0–12): Boston density. Chains seeded from published disclosures (the Sweetgreen / Halal Guys pattern, already demonstrated); independents digitized block-by-block under a documented collection protocol (source → digitize → estimate → owner-verify → maintain). Acquisition via TikTok/influencer (the Cal AI playbook) aimed at GLP-1 and goal-oriented audiences — trust-first, guilt-free positioning the incumbents ceded.

The pilot is pre-registered. Six metrics defined *before* measurement, published on the product's public Impact page: pre-order decision rate; verified-menu coverage; correction turnaround; logged-meal accuracy (order ground truth vs. estimates); independent-restaurant onboarding cost; diet-quality delta. Definitions don't move after the fact.

Phase 2 (12–30 mo): convert coverage into retail-media; raise seed on retention + featured-restaurant proof (AI took 62% of 2025 digital-health VC). **Phase 3 (30 mo+):** payer/employer contracts on covered MNT.

8 Funding strategy

- **Pre-seed:** ~\$1.0–1.5M SAFE (~\$5–6M post) — mid-range of 2024–26 norms, defensible with a working product and pilot protocol.
- **Non-dilutive:** NIH/NIDDK, USDA-NIFA, NSF SBIR (Phase I ~\$150–300k / \$125–175k / ~\$275k; Phase II \$1–2M). *SBIR/STTR lapsed Oct 2025, reauthorized Apr 2026 — verify solicitation calendars.*
- **Seed (~\$4–6M) at months 12–18** on retention + featured-cohort proof; **Series A (~\$15–25M)** for GTM and the payer leg.
- **Why now:** the GLP-1 / food-as-medicine funding wave (Nourish \$70M Series B at \$1B+; Foodsmart \$200M; Fay and Berry Street \$50M each) rewards evidence-first nutrition positioning.

9 Milestones — next 18 months

Milestone	Target	Status
Working end-to-end product (decision → order → log → measure)	—	Done, Jul 2026
Published-nutrition chain seeding demonstrated	—	Done
Pilot metrics pre-registered publicly	—	Done
Backend: cross-device accounts, live terminal sync	Mo 1–2	Next
Catalog: 20+ restaurants incl. verified independents	Mo 3	13 today
First 3–5 owner-verified partners (letters + correction logs)	Mo 4–6	—
D30 retention > 35% (vs. ~30% category)	Mo 6–12	—
Featured-restaurant cohort 150+ · ~3,000 Boston subscribers	Mo 12–18	—
SBIR Phase I · Seed raised	Mo 9–18	—

10 Risks & mitigations

Risk	Mitigation
Estimation accuracy is physics-bound (~16–25% photo floor)	Lead with menu text (models match nutritionists there); ground in USDA FNDDS (grounding cuts error 63–83%); surface $\pm$ ranges; correction built into every surface; never marketed as clinical
Retention kills this category	Plan-ahead is a decision habit, not a logging chore; order-confirmed logging removes the diary; calibration compounds the reason to stay; annual plans ~halve churn
Two-sided cold start	Published-disclosure seeding gives day-one chain coverage; density beats breadth; the terminal gives independents value before any fee
GLP-1 substitution	Position as the eating-out companion <i>for</i> GLP-1 users — the angle funding the category
Incumbent consolidation (MFP $\times$ Cal AI)	Moat is the verified-data + correction + ground-truth graph, not the model
Trust/integrity failure would be fatal	Governance is product: sponsored-scoring separation is categorical; corrections versioned and public; simulated states labeled; a standing no-unrecorded-claims commitment
Policy risk on B2B2C	Anchor on commercial MNT; Medicaid expansion as optionality

11 National-interest alignment

This section states the founders' case factually; it draws no legal conclusions.

- **Substantial merit.** Restaurant-nutrition transparency for the independent majority federal labeling does not reach; pre-order decision support where labels demonstrably fail (~24 kcal/order); health literacy through explained, sourced, correctable numbers; small-restaurant menu digitization at documented cost.
- **National importance.** A reproducible protocol — collection standard, provenance tiers, correction governance, pre-registered measurement — designed to scale across U.S. markets; addressed at a diet-disease burden of ~\$334B/yr.
- **Well positioned.** Working product with dated commit history; live-captured published-data integrations; documented methods; financial and competitive analyses; a pilot designed for verifiable evidence (restaurant letters, correction logs, exportable provenance-carrying data).
- **Evidence integrity.** No fabricated users, restaurants, partners, outcomes, revenue, ratings, or health claims — enforced in the product itself (prototype stamps, simulation labels, demo-payment disclosures, confidence tiers).

12 Team

Founding team — two co-founders: one leading product and engineering (builder of the product to date), one leading strategy, market validation, and regulatory-economic analysis. Founder credentials are shared with investors directly. Pre-seed hires: a registered dietitian / nutrition-science advisor (credibility + B2B2C), an

ML engineer (estimation engine), and a local GTM lead (restaurant density). Advisory targets: a fast-casual operator and a digital-health payer-contracting advisor.

Appendices available on request: Financial Model · Market Research Dossier (cited evidence base) · Competitive Scan, July 2026 (live captures) · Menu Collection Protocol · System Architecture. A live product demonstration accompanies this plan.