

---

# Forkcast

## The trust layer for restaurant nutrition decisions

Business Plan — Version 2.0 · July 2026

Prepared for pre-seed conversations and the founder's evidence record.

---

### THE PRODUCT IN ONE SENTENCE

Forkcast tells diners what to order *before* they order — scored against their own body and goals — then closes the loop: order, confirm, log, measure. It is the only product connecting pre-order decision support, verified independent-restaurant nutrition data, and post-order confirmed logging.

### SOURCING & INTEGRITY

All market figures are sourced (CDC, USDA ERS, AHRQ, peer-reviewed literature, company filings/press; see `RESEARCH.md`). Competitive claims come from live site captures dated July 15, 2026 (`COMPETITIVE_SCAN_2026-07.md`). Nothing in this document claims users, partners, revenue, or outcomes that do not exist.

Working product: private repository `Seymurhh/forkcast` (38+ dated commits) · live demo site

Companions: `FINANCIAL_MODEL.md` · `RESEARCH.md` · `MENU_COLLECTION.md` · `SYSTEM_ARCHITECTURE.md`

## 0 The answer first

**Americans eat out more than ever, decide blind, and every existing app reacts only after the plate is empty.** The category leader for “eat out healthy” markets AI guesses across 22 million locations as facts; the logging giants own the aftermath; the delivery marketplaces own the transaction. Nobody owns the *decision* — and nobody at all serves the roughly 65% of U.S. restaurants that are independent and exempt from federal menu labeling.

**Forkcast owns the decision, and earns the right to with honesty no competitor structurally can match:** every nutrition number carries its source and a  $\pm$  range; restaurants review and correct their own data through a versioned public process; sponsored placement can never touch the score; and the personal calorie target itself recalibrates from the user’s own measured energy balance. The full product — decision → order → restaurant terminal → confirmed log → adaptive measurement — **exists and works today.**

**The ask:** ~\$1.0–1.5M pre-seed SAFE (~\$5–6M post) plus non-dilutive SBIR, to run the Boston pilot against six pre-registered metrics, convert the first independent restaurants into verified partners, and prove retention economics.

## 1 Situation — eating out is where American diets are decided

| Fact                                          | Figure                       | Source                                            |
|-----------------------------------------------|------------------------------|---------------------------------------------------|
| Food spending eaten away from home            | <b>58.5%</b>                 | USDA ERS, 2024 (record; \$4,485/capita, +12% YoY) |
| Daily calories from food away from home       | ~1/3                         | USDA ERS / NHANES (directional)                   |
| Restaurant vs. home meal                      | <b>+200 kcal, +350 mg Na</b> | USDA ERS                                          |
| Diners underestimating restaurant calories    | <b>2 in 3</b>                | USDA/NHANES 2024 (~25% by 500+ kcal)              |
| Behavior change from mandatory calorie labels | <b>~24 kcal/order</b>        | Peer-reviewed labeling literature                 |
| U.S. adults obese / overweight                | <b>40.3% + 31.7%</b>         | CDC NCHS, 2024 (~72% combined)                    |
| Diet-related disease cost                     | <b>~\$334B/yr</b>            | CDC; AHRQ MEPS (obesity ~\$173B, 2019\$)          |


## 2 Complication — three structural failures nobody has fixed

- The data doesn’t exist where people actually eat.** FDA menu labeling (21 CFR 101.11) binds only 20+-location chains. The independent majority publish nothing, and no commercial database covers them.
- The timing is wrong everywhere.** Logging apps (MyFitnessPal: 220M registered, ~30M MAU) act after eating — accounting, not coaching — and 70–80% of calorie-app users quit within two weeks. Labels act at the menu but move behavior ~24 kcal. Delivery apps act at the transaction with zero nutrition intelligence.

3. **The new “AI menu” entrants sell certainty they don’t have.** MenuFit (live capture, July 2026): “#1 App For Eating Out Healthy,” 22.3M locations, meals that “perfectly align with your macros” — no sources, no error ranges, no corrections, no restaurant participation, no web product, and a recommendation that dead-ends. Credibility is the category’s open flank.

### 3 Resolution — the trust layer, already built

Forkcast turns a health profile (Mifflin-St Jeor BMR → activity-scaled TDEE → goal-adjusted targets; ISSN macros; CDC BMI classes; allergies and conditions) into an explainable per-dish **Fit Score**, ranks every nearby menu around what’s left of the user’s day, and then — unlike every competitor — finishes the job:




**Exhibit 1 — The closed loop.** Competitors stop at the recommendation (or start after the meal); Forkcast connects the decision to the record. Highlighted steps are where every competitor drops out.

#### As built — dated commits, July 2026

| Capability (all live)                                               | Why it matters                                                                                                                                                                           |
|---------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Fit Score engine + condition/allergen personalization               | Explainable, safety-aware ranking; allergen matches excluded from recommendations; condition advisories on every dish                                                                    |
| Full ordering loop: basket → checkout → tracking → “Log this meal?” | The decision connects to the record; every logged restaurant meal carries an order reference                                                                                             |
| Restaurant partner terminal                                         | Order queue with customer allergy flags, prep quotes, and <b>versioned, timestamped menu corrections</b> surfacing publicly on dish pages                                                |
| Three-tier data provenance                                          | <i>Partner-verified</i> / <i>Restaurant-published</i> (Sweetgreen captured live July 2026; The Halal Guys’ published guide) / <i>Estimated ± 15%</i> (real Boston independents, labeled) |
| Adaptive metabolic calibration                                      | Target TDEE back-calculated from the user’s own logged intake + weight change, confidence-blended — the state-of-the-art approach, attached to restaurant decisions                      |
| Evidence system                                                     | Exportable meal log with source/portion/confidence per entry; / <i>impact</i> page with six pre-registered pilot metrics and a public source ledger                                      |
| Role-based accounts; web-first, shareable dish pages                | Two-sided from day one; app-only competitors have no web surface at all                                                                                                                  |

**Deliberately *not* built:** payment processing, live delivery dispatch, partner handoff links — all shown as clearly-labeled integration states, never simulated as real.

## System architecture




**Exhibit 2 — Architecture.** Three surfaces over shared engines over the data asset. Every arrow that touches the data layer makes it more defensible: corrections from restaurants, order-confirmed ground truth from diners. Full detail: `SYSTEM_ARCHITECTURE.md`.

## 4 Market

U.S.-only; reasoning explicit so each number can be challenged (derivations in `RESEARCH.md`).


- **TAM** ≈ **\$11.1B/yr.** ~186M U.S. adults with elevated BMI (≈72% of 258M; CDC 2024) × ~\$60 blended annual ARPU (anchored between MyFitnessPal Premium \$79.99/yr and Cal AI \$29.99/yr). Triangulates with the independently sized calorie-tracking app market (\$5.5–14B, 2025; 12–20% CAGR).
- **SAM** ≈ **\$3.3B/yr.** ~55.7M goal-oriented active dieters who eat out often (~30% of the pool) × \$60.
- **SOM** ≈ **\$40–90M ARR (5-yr).** 0.5–2.0% of SAM (0.28–1.1M payers; 1% ≈ \$33M consumer ARR) plus restaurant retail-media and B2B2C legs. **Retention-bound, not demand-bound.**
- **Velocity proof:** Cal AI reached ~15M downloads and \$30M+ ARR in under two years before MyFitnessPal acquired it (Dec 2025).



**Exhibit 3 — Market pyramid** (widths to scale within each tier's share of the tier above).

## 5 Competition — and the flank they all leave open

| Player                      | Owns                                                  | Structurally absent                                                                                                                                         |
|-----------------------------|-------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------|
| MenuFit / PlateMate / Nuuro | “Know what to order”; AI menu scoring; influencer GTM | Sources, error ranges, corrections, restaurant participation, ordering/logging loop, web presence; guilt-based framing excludes the health-condition market |
| MyFitnessPal                | Retrospective logging at scale                        | Acts after eating; restaurant data weakest (same dish 400–1,200 kcal); 70–80% two-week churn; first revenue decline (–5.7%, 2025)                           |
| DoorDash / Uber Eats        | Ordering logistics, merchant graph, \$1B+ ads         | Zero personal-nutrition intelligence; sponsored placement <i>is</i> the product                                                                             |
| Sweetgreen / CAVA           | Excellent single-brand transparency                   | Cannot be brand-agnostic                                                                                                                                    |
| Foodsmart / Nourish         | Payer-funded dietitian care                           | Care delivery, not the moment of ordering                                                                                                                   |



**Exhibit 4 — Positioning.** The upper-right quadrant — decide before ordering, on data you can prove — is empty. Chains publish but serve only themselves; the AI entrants act early but cannot prove a number; the giants act late or not at all.

**Whitespace in one line:** *nobody connects pre-order decision support, verified independent-restaurant data, and post-order confirmed logging — and nobody can prove their numbers.*

**The moat compounds through participation:** every verified menu, every versioned correction, every order-confirmed meal (ground truth vs. estimate — a validation dataset no retrospective logger can construct), and every calibration cycle makes the data asset harder to replicate than any vision model.

## 6 Business model — three legs, sequenced

1. **Consumer subscription (launch) — priced to win.** The closed loop is **free forever** (charging for logging is the incumbents’ most-hated move, and free order volume feeds the restaurant leg). Premium — calibration, analytics, unlimited photo AI, condition packs — at **\$4.99/mo or \$39.99/yr**: half of MyFitnessPal’s annual,

- 4–14× under every incumbent monthly. Gate growth on **LTV:CAC**  $\geq 3:1$ , ~6-month payback, funded by a ~\$0-CAC channel no logger has: restaurant-partner co-marketing.
2. **Restaurant leg — every dollar earned (incentive ladder).** Rung 1: **Verified Partner is free forever** (terminal, allergy flags, corrections, badge) — their participation improves our data asset. Rung 2: **6% commission on Forkcast-originated pickup orders** — against the 15–30% they already pay marketplaces; they pay only when we deliver a customer. Rung 3: optional **featured placement from \$99/mo**, bought after the order channel proves volume — and **never affecting Fit Scores**. Blended yield ~\$1,050/location/yr.
  3. **B2B2C covered nutrition (durable).** Payer/employer contracts on covered MNT (CPT 97802–97804); GLP-1 companion programs; Medicaid expansion as optionality.

### The price war, won on entry

| Product                 | Monthly       | Annual         | Note                                                              |
|-------------------------|---------------|----------------|-------------------------------------------------------------------|
| MyFitnessPal Premium    | \$19.99       | \$79.99        | Premium+ \$24.99/\$99.99; paywall backlash, first revenue decline |
| Noom / coaching apps    | \$40–70       | ~\$209         | \$62M dark-pattern settlement                                     |
| MacroFactor             | \$11.99       | \$71.99        | Monetizes accuracy/adaptivity — our closest analog                |
| Cal AI                  | —             | \$29.99        | Photo-only volume play                                            |
| <b>Forkcast Premium</b> | <b>\$4.99</b> | <b>\$39.99</b> | The closed loop itself stays free forever                         |

**Unit economics at the new price (v2):** blended gross ARPU ~\$42/yr; CAC ~\$15 (restaurant co-marketing as a ~\$0-CAC channel); **LTV:CAC**  $\approx 3.3:1$ , **payback**  $\approx 6$  months. The honest trade: Year-5 base ARR moves from \$58M (v1) to a more defensible **\$37M** — low end of the \$30–90M SOM envelope — with named upside paths (featured adoption, order volume, earned ARPU lifts, GLP-1 employer channel) back toward \$50M+. Pricing headroom is *upward*: verified trust can raise prices; eroded trust can only cut. Full build: FINANCIAL\_MODEL.md (v2).

## 7 Go-to-market — density, then proof, then scale

**Phase 1 (months 0–12): Boston density.** Chains seeded from published disclosures (the Sweetgreen / Halal Guys pattern, already demonstrated); independents digitized block-by-block under the documented protocol (MENU\_COLLECTION.md: source → digitize → estimate → owner-verify → maintain). Acquisition via TikTok/influencer (the Cal AI playbook) aimed at GLP-1 and goal-oriented audiences — trust-first, guilt-free positioning the incumbents ceded.

**The pilot is pre-registered.** Six metrics defined *before* measurement, published at /impact: pre-order decision rate; verified-menu coverage; correction turnaround; logged-meal accuracy (order ground truth vs. estimates); independent-restaurant onboarding cost; diet-quality delta. Definitions don't move after the fact.

**Phase 2 (12–30 mo):** convert coverage into retail-media; raise seed on retention + featured-restaurant proof (AI took 62% of 2025 digital-health VC). **Phase 3 (30 mo+):** payer/employer contracts on covered MNT.

## 8 Funding strategy

- **Pre-seed:** ~\$1.0–1.5M SAFE (~\$5–6M post) — mid-range of 2024–26 norms, defensible with a working product and pilot protocol.

- **Non-dilutive:** NIH/NIDDK, USDA-NIFA, NSF SBIR (Phase I ~\$150–300k / \$125–175k / ~\$275k; Phase II \$1–2M). *SBIR/STTR lapsed Oct 2025, reauthorized Apr 2026 — verify solicitation calendars.*
- **Seed (~\$4–6M) at months 12–18** on retention + featured-cohort proof; **Series A (~\$15–25M)** for GTM and the payer leg.
- **Why now:** the GLP-1 / food-as-medicine funding wave (Nourish \$70M Series B at \$1B+; Foodsmart \$200M; Fay and Berry Street \$50M each) rewards evidence-first nutrition positioning.

## 9 Milestones — next 18 months

| Milestone                                                     | Target   | Status                |
|---------------------------------------------------------------|----------|-----------------------|
| Working end-to-end product (decision → order → log → measure) | —        | <b>Done, Jul 2026</b> |
| Published-nutrition chain seeding demonstrated                | —        | <b>Done</b>           |
| Pilot metrics pre-registered publicly (/impact)               | —        | <b>Done</b>           |
| Backend: cross-device accounts, live terminal sync            | Mo 1–2   | Next                  |
| Catalog: 20+ restaurants incl. verified independents          | Mo 3     | 13 today              |
| First 3–5 owner-verified partners (letters + correction logs) | Mo 4–6   | —                     |
| D30 retention > 35% (vs. ~30% category)                       | Mo 6–12  | —                     |
| Featured-restaurant cohort 150+ · ~3,000 Boston subscribers   | Mo 12–18 | —                     |
| SBIR Phase I · Seed raised                                    | Mo 9–18  | —                     |

## 10 Risks & mitigations

| Risk                                                       | Mitigation                                                                                                                                                                                        |
|------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Estimation accuracy is physics-bound (~16–25% photo floor) | Lead with menu text (models match nutritionists there); ground in USDA FNDDS (grounding cuts error 63–83%); surface $\pm$ ranges; correction built into every surface; never marketed as clinical |
| Retention kills this category                              | Plan-ahead is a decision habit, not a logging chore; order-confirmed logging removes the diary; calibration compounds the reason to stay; annual plans ~halve churn                               |
| Two-sided cold start                                       | Published-disclosure seeding gives day-one chain coverage; density beats breadth; the terminal gives independents value before any fee                                                            |
| GLP-1 substitution                                         | Position as the eating-out companion <i>for</i> GLP-1 users — the angle funding the category                                                                                                      |
| Incumbent consolidation (MFP × Cal AI)                     | Moat is the verified-data + correction + ground-truth graph, not the model                                                                                                                        |
| Trust/integrity failure would be fatal                     | Governance is product: sponsored-scoring separation is categorical; corrections versioned and public; simulated states labeled; /impact no-unrecorded-claims commitment                           |
| Policy risk on B2B2C                                       | Anchor on commercial MNT; Medicaid expansion as optionality                                                                                                                                       |

## 11 National-interest alignment

---

This section states the founder's case factually; it draws no legal conclusions.

- **Substantial merit.** Restaurant-nutrition transparency for the independent majority federal labeling does not reach; pre-order decision support where labels demonstrably fail (~24 kcal/order); health literacy through explained, sourced, correctable numbers; small-restaurant menu digitization at documented cost.
- **National importance.** A reproducible protocol — collection standard, provenance tiers, correction governance, pre-registered measurement — designed to scale across U.S. markets; addressed at a diet-disease burden of ~\$334B/yr.
- **Well positioned.** Working product with dated commit history; live-captured published-data integrations; documented methods; financial and competitive analyses; a pilot designed for verifiable evidence (restaurant letters, correction logs, exportable provenance-carrying data).
- **Evidence integrity.** No fabricated users, restaurants, partners, outcomes, revenue, ratings, or health claims — enforced in the product itself (prototype stamps, simulation labels, demo-payment disclosures, confidence tiers).

## 12 Team

---

Founder: Harvard SEAS (engineering); sole builder of the product to date. Pre-seed hires: a registered dietitian / nutrition-science advisor (credibility + B2B2C), an ML engineer (estimation engine), and a local GTM lead (restaurant density). Advisory targets: a fast-casual operator and a digital-health payer-contracting advisor.

---

Companions: `FINANCIAL_MODEL.md` · `RESEARCH.md` (cited evidence base) · `COMPETITIVE_SCAN_2026-07.md` (live captures) · `MENU_COLLECTION.md` (data protocol) · `SYSTEM_ARCHITECTURE.md` · `DESIGN_INTEGRATION_NIW.md`. Working product: repository (`npm run dev`) and the published demo site.